




## 1. GOVERNING PURPOSE

The Agentic AI Assurance Workforce Cone Framework™ establishes a human assurance-capability architecture for governing Agentic AI systems throughout their lifecycle.

It addresses:

**COMPETENCE → INDEPENDENCE → AUTHORITY → ASSURANCE → MONITORING → EVIDENCE  
→ INTERVENTION → REVALIDATION → REVOCATION**

The Framework is designed to ensure that consequential Agentic AI capability is not treated as deployment-ready merely because the underlying technology is technically capable.



A nominal appointment without appropriate competence, authority, access, independence where required, resources, continuity, or required coverage shall not constitute effective assurance readiness. Assurance readiness shall be evidenced and shall remain effective for the applicable risk, operating conditions, and authority granted.

This gate applies proportionately to the system's risk, autonomy, authority, connectivity, complexity, and potential consequence.


Where applicable law, regulation, contractual obligation, professional requirement, or binding supervisory requirement imposes a stricter condition, that requirement shall prevail.

### 3. ASSURANCE WORKFORCE MANDATE

The Assurance Workforce exists to provide appropriately qualified human capability to:

**DEFINE → ASSESS → TEST → AUTHORIZE → DEPLOY → MONITOR → EVIDENCE → INTERVENE  
→ INVESTIGATE → REVALIDATE → REAUTHORIZE / RESTRICT / REVOKE**

The required level of assurance shall be proportionate to:



Purely informational, read-only, advisory, or otherwise non-consequential use may be subject to a proportionately reduced assurance posture, unless a stricter applicable requirement applies.

**CONTROL-BOUNDARY PRINCIPLE** — This Framework governs assurance readiness and control conditions; it does not by itself determine whether a particular AI use is legally permitted, prohibited, licensed, or subject to regulatory approval. Those determinations remain subject to the applicable jurisdiction, sector, activity, and competent authority.

## 5. RISK TIERING & ASSURANCE CADENCE REFERENCE

The Framework calibrates competency reassessment, monitoring intensity, evidence retention, and revalidation to the applicable risk tier. Tier assignment shall be documented, justified, approved by the designated governance authority, and reassessed when material changes alter the risk profile.

The stated cadences are illustrative baseline minimums for organizational adoption, not universal legal or regulatory requirements.

Organizations shall document the applicable tier as part of Pre-Deployment Assurance and shall apply any stricter mandatory legal, regulatory, contractual, professional, or organizational requirement.

### Tier 1 — Critical

Examples include autonomous financial transactions above materiality thresholds, safety-affecting actions, privileged/security-relevant access, irreversible external commitments, and regulated-sector consequential decisions.

- Competency reassessment: every 6 months and immediately upon material change
- Monitoring: real-time continuous automated monitoring with immediate alerting
- Evidence retention: 7 years, or applicable regulatory period if longer
- Revalidation: every 90 days and immediately upon material change

### Tier 2 — High

Examples include consequential customer- or employee-facing decisions, sensitive/confidential data processing, delegated authority to other systems, and external interaction with material business effect.

- Competency reassessment: every 12 months
- Monitoring: near-real-time alerting within 15 minutes
- Evidence retention: 3 years, or applicable regulatory period if longer
- Revalidation: every 180 days

### Tier 3 — Standard


Examples include internal operational actions with limited external effect and reversible transactions within pre-approved limits.

- Competency reassessment: every 12-18 months
- Monitoring: daily aggregated-log review
- Evidence retention: minimum 1 year
- Revalidation: annually

### Tier 4 — Low

Examples include read-only, advisory, informational, or otherwise non-consequential actions.

- Competency reassessment: every 24 months
- Monitoring: periodic sampling, minimum weekly
- Evidence retention: minimum 1 year
- Revalidation: annually or upon material change




## THESE ARE BASELINE MINIMUMS, NOT CEILINGS.

Organizations may adopt shorter cadences, tighter thresholds, stronger monitoring, or more frequent revalidation. A stated baseline does not create permission to operate where a stricter mandatory requirement, unresolved material control deficiency, or loss of required assurance capability prevents continued authorization.

Nothing in this Section limits a more conservative control posture.

## 6. ASSURANCE WORKFORCE ESTABLISHMENT

Before consequential deployment, the organization shall establish, as appropriate to the system's risk and operating environment:



## 7. DEFINITION OF QUALIFIED & AUTHORIZED ASSURANCE WORKFORCE

For purposes of this Framework, a Qualified & Authorized Assurance Workforce means designated human personnel or an appropriately constituted human assurance function that:

1. Possesses demonstrable competence appropriate to the system's risk and operating environment;
2. Has been formally assigned assurance responsibilities;
3. Has sufficient authority and access to perform those responsibilities;
4. Has appropriate independence where required by law, regulation, professional requirements, governance arrangements, or control design;

5. Has adequate resources, coverage, continuity, and escalation capability;
6. Can review relevant evidence and challenge material outcomes;
7. Can exercise required intervention, restriction, suspension, escalation, or revocation authority; and
8. Is subject to periodic competency and authorization revalidation.

A job title, nominal appointment, training attendance, technical access, or organizational designation alone shall not establish qualification or effective assurance readiness.

## 8. COMPETENCY REQUIREMENT

Personnel shall possess demonstrable competence appropriate to:

**RISK → AUTONOMY → AUTHORITY → COMPLEXITY → CONSEQUENCE → OPERATING ENVIRONMENT**

An Agentic AI system can be technically capable — but capability alone does **NOT** make it

Relevant competency domains may include:

**AGENTIC AI GOVERNANCE → AI RISK → INTERNAL AUDIT & ASSURANCE → CYBERSECURITY → TECHNOLOGY RISK → SYSTEM / MODEL EVALUATION → DATA GOVERNANCE → PRIVACY → REGULATORY COMPLIANCE → AUTHORITY MANAGEMENT → LOGGING & EVIDENCE → INCIDENT RESPONSE → FORENSICS → CONTAINMENT → REVOCATION & RECOVERY**

NO ASSURANCE READINESS

Competence may be demonstrated through:

**QUALIFICATIONS → CERTIFICATIONS → TRAINING → EXPERIENCE → ASSESSMENT → DEMONSTRABLE PRACTICAL COMPETENCY**

### THE ASSURANCE ARCHITECTURE

Where applicable law, regulation, professional requirements, licensing conditions, accreditation requirements, or binding organizational requirements prescribe a specific credential or qualification, that requirement shall be satisfied.

**THIS FRAMEWORK DOES NOT ITSELF CREATE, IMPLY, OR REPRESENT A GOVERNMENT-ISSUED CERTIFICATION, LICENSE, ACCREDITATION, OR STATUTORY QUALIFICATION.**

Hanan Shamim  
Founder of The First Agentic AI Audit Ecosystem™

## 9. INDEPENDENCE & CONFLICT MANAGEMENT

Where independent assurance is required by applicable law, regulation, professional requirements, governance arrangements, or adopted control design, the assurance function shall have appropriate:


**INDEPENDENCE → AUTHORITY → ACCESS → COMPETENCE → RESOURCES**

Material conflicts shall be:

**IDENTIFIED → DISCLOSED → ASSESSED → MITIGATED → DOCUMENTED → ESCALATED WHERE REQUIRED**

No operational party shall be permitted to unilaterally:

- Suppress material findings
- Alter material assurance evidence
- Destroy required evidence
- Redefine a material failure as success
- Disable mandatory controls
- Prevent required escalation
- Circumvent required assurance
- Override an assurance conclusion outside the applicable governance process



## 11. AUTHORIZATION SHALL BE EXPLICIT


TECHNICAL CAPABILITY ≠ AUTHORIZATION

Authorization shall be:

**EXPLICIT → ATTRIBUTABLE → SCOPED → CONDITIONED → CURRENT → EVIDENCED**

The following shall not, by themselves, constitute authorization:

- Technical readiness
- Successful testing
- Vendor approval
- User acceptance
- Management intention
- Historical safe operation
- Absence of incidents



**AI CAPABILITY SHALL NEVER MANUFACTURE HUMAN AUTHORITY.**

### 13. HUMAN OVERSIGHT MUST BE EFFECTIVE

Where human intervention is required by risk, law, regulation, governance, or control design, authorized personnel shall have sufficient capability and authority to:

**UNDERSTAND → ASSESS → CHALLENGE → REJECT / OVERRIDE WHERE FEASIBLE → RESTRICT → ESCALATE → SUSPEND → SAFE-STOP → REVOKE**

Observation without meaningful intervention capability shall not be treated as effective oversight where intervention is required.

Human presence alone does not establish effective human oversight.

## 14. CONTINUOUS ASSURANCE

Post-deployment assurance shall be proportionate to risk and shall meet, at minimum, the applicable monitoring thresholds established under Section 5, subject to stricter applicable requirements.

Monitoring shall address, as applicable:



controls shall itself constitute a control event.

Where required by risk:

**RESTRICT → SUSPEND → SAFE-STATE**

Continued technical operation shall not, by itself, constitute evidence of governance effectiveness.

## 16. MATERIAL EVENT RESPONSE

DETECT → ALERT → ESCALATE → CONTAIN → RESTRICT / SUSPEND → PRESERVE EVIDENCE → INVESTIGATE → REMEDIATE → REVALIDATE → REAUTHORIZE / REVOKE

The response shall be appropriately documented and evidence-bounded. Material decisions, interventions, exceptions, and disposition outcomes shall be attributable to the responsible human authority or authorized function.

## 17. DYNAMIC AUTHORIZATION

Authorization shall not be treated as permanently valid. Authorization shall automatically cease, restrict, or require reassessment when its defined scope, conditions, duration, authority, environment, version, tools, data, delegation, or other mandatory authorization condition is no longer satisfied.

Authorization shall remain bounded by:

SCOPE → CONDITIONS → DURATION → AUTHORITY → ENVIRONMENT → VERSION → TOOLS → DATA → DELEGATION

An Agentic AI system can be technically capable — but capability alone does **NOT** make it deployment-ready.

Material changes to the model, agent, tools, permissions, credentials, objectives, workflows, data, environment, delegation, configuration, legal/regulatory obligations, or risk profile shall trigger reassessment before affected consequential authority is continued or expanded, unless an applicable emergency-control process expressly permits temporary operation within defined limits.

Absent a more specific trigger, reassessment shall occur at least as frequently as the applicable revalidation cadence in Section 5.

## 18. REVOCATION MUST BE DEMONSTRABLE

Revocation shall have a defined technical and operational control path.

The organization shall be able, where relevant, to establish:

WHO REVOKED → WHAT WAS REVOKED → WHEN → WHY → CONTROL PATH → EFFECTIVE STATUS → EXECUTION  
STATUS → EVIDENCE

**AI WORKS.**  
Technology delivers capability.

**HUMANS DECIDE.**  
People set boundaries and intent.

**CONTROLS ENFORCE.**  
Systems and governance ensure compliance.

**ASSURANCE PROVES.**  
Evidence builds trust and confidence.

**REVOCATION REMAINS POSSIBLE.**  
Control is never permanent.

**A STOP INSTRUCTION IS NOT, BY ITSELF, PROOF THAT EXECUTION STOPPED.**

*Adnan Shamim*

Founder of The First Agentic AI Audit Ecosystem™

Where technically or operationally feasible, revocation effectiveness shall be independently verifiable.

Revocation evidence shall be retained for no less than the applicable minimum period defined by risk tier, or any longer applicable mandatory retention period.

## 19. EMERGENCY CONTROL — NO PERMANENT BACKDOOR

Exceptional or degraded-control operation shall be:

EXPLICIT → SCOPE-BOUND → TIME-BOUND → ATTRIBUTABLE → LOGGED → MONITORED → REVIEWED → EXPIRED / REVOKED

Emergency mechanisms shall not become permanent uncontrolled authorization pathways.



## 22. ASSURANCE LIFECYCLE

### DEFINE

Purpose, boundaries, risk, accountability, and intended authority.




### ASSESS

Architecture, risks, controls, workforce capability, and operating environment.



### TEST

Material controls and consequential execution paths.



### INVESTIGATE

Determine what occurred and what evidence supports the conclusion.



### REVALIDATE

Reassess continued authorization and control effectiveness.



### REAUTHORIZE / RESTRICT / REVOKE

Continue, limit, suspend, or terminate authority.

## 23. FINAL HARD CONTROL

The pre-deployment assurance gate established in Section 2 remains applicable throughout the operational life of a consequential Agentic AI system.

Post-deployment, it extends further:

### NO REQUIRED OVERSIGHT / MONITORING



### NO UNRESTRICTED CONTINUED OPERATION

Where required by applicable risk, governance, law, regulation, or control design, loss of required assurance capability shall trigger proportionate restriction, suspension, safe-state operation, or other defined intervention.



### EVIDENCE



### INTERVENTION



### REVALIDATION



### REVOCATION


## 26. LEGAL, REGULATORY & ADOPTION SAFEGUARD

### FOR INFORMATIONAL AND INTERNAL GOVERNANCE PURPOSES ONLY

This Framework is provided for informational, educational, and internal governance purposes only.

It is a conceptual governance and assurance architecture and does not constitute:

- Legal advice
- Regulatory advice
- A legal opinion



The Framework does not itself create or imply any:

**GOVERNMENT CERTIFICATION → STATUTORY QUALIFICATION → REGULATORY PERMISSION → PROFESSIONAL LICENSE → ACCREDITATION → LEGAL AUTHORITY**

## 27. FINAL CONTROL PRINCIPLES — QUICK-REFERENCE INDEX

AI MAY ACT ONLY WITHIN AUTHORIZED BOUNDARIES.

HUMANS RETAIN ACCOUNTABILITY.

ASSURANCE MUST BE INDEPENDENT WHERE REQUIRED.

THE AGENTIC AI ASSURANCE WORKFORCE CONE FRAMEWORK™

**AGENTIC AI DOESN'T JUST  
NEED AUTONOMY.  
IT NEEDS ASSURANCE.**

**MATERIAL CONTROLS MUST BE TESTABLE.**

**MATERIAL ACTIONS MUST BE GOVERNABLE.**

but capability alone does **NOT** make it deployment-ready.

**MATERIAL CLAIMS MUST BE EVIDENCE-BOUNDED.**

**MATERIAL AUTHORITY MUST BE REVOCABLE.**

**MONITORING FAILURE IS ITSELF A CONTROL EVENT.**

**AUTHORIZATION IS DYNAMIC, NOT PERPETUAL.**

**AI CANNOT SELF-AUTHORIZE.**

**AI CAPABILITY DOES NOT CREATE HUMAN AUTHORITY.**

**REGULATORY AUTHORITY CANNOT BE SELF-GRANTED BY AI.**


**CONTROLLED AUTONOMY.  
DEMONSTRABLE ASSURANCE.  
REAL ACCOUNTABILITY.**

The future of Agentic AI is not simply about increasing autonomy. It is about making autonomy **governable, controlled, observable, testable, evidence-bounded, and revocable.**  
**Prove it. Don't just state it.**



THE ASSURANCE CHAIN  
**I-A-N-E-T-F™**

- I** Intent
- A** Authority
- N** Enforcement
- E** Testing
- T** Traceability & Evidence
- F** Forensics & Revocability



*Adnan Shamim*

Founder of The First Agentic AI Audit Ecosystem™

LEGAL AND REGULATORY REQUIREMENTS PREVAIL WHERE APPLICABLE.

ASSURANCE READINESS MUST EXIST BEFORE CONSEQUENTIAL AUTONOMY.

**FINAL ASSURANCE BOUNDARY:** A framework assertion is not evidence of implementation. Control effectiveness, authorization status, compliance, and revocation effectiveness must be established from the applicable implementation's own authoritative evidence.

